Network architecture: function NNθ: R

input−dim → R

output−dim parametrized by θ ≡ (W, b)

+

Loss function L(θ): is minimized during training.

+

Optimization procedure (gradient descent): updates θ with ∇L(θ); θ ← θ − η∇L(θ)